

```
import nltk
nltk.download('stopwords')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.
True

import pandas as pd
import numpy as np
import re
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import style
style.use('ggplot')
from nltk.tokenize import word_tokenize
from nltk.stem import WordNetLemmatizer
from nltk.corpus import stopwords
stop_words = set(stopwords.words('english'))
from wordcloud import WordCloud
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, ConfusionMatrixDisplay

from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

```
tweet_df = pd.read_csv("/content/drive/MyDrive/train.csv")
tweet_df.head()
```

	id	label	tweet
0	1	0	@user when a father is dysfunctional and is s...
1	2	0	@user @user thanks for #lyft credit i can't us...
2	3	0	bihday your majesty
3	4	0	#model i love u take with u all the time in ...
4	5	0	factsguide: society now #motivation

Next steps: [Generate code with tweet_df](#) [View recommended plots](#)

```
tweet_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31962 entries, 0 to 31961
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0    id      31962 non-null    int64
```


True

```
tweet_df['tweet'] = tweet_df['tweet'].apply(lambda x: lemmatizing(x))
```

```
# printing the data to see the effect of preprocessing
```

```
print(tweet_df['tweet'].iloc[0], "\n")
```

```
print(tweet_df['tweet'].iloc[1], "\n")
```

```
print(tweet_df['tweet'].iloc[2], "\n")
```

```
print(tweet_df['tweet'].iloc[3], "\n")
```

```
print(tweet_df['tweet'].iloc[4], "\n")
```

```
user father dysfunctional selfish drags kids dysfunction run
```

```
user user thanks lyft credit cant use cause dont offer wheelchair vans pdx disapointed getthanked
```

```
bihday majesty
```

```
model love u take u time ur
```

```
factsguide society motivation
```

```
tweet_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 29345 entries, 0 to 31961
```

```
Data columns (total 3 columns):
```

```
#   Column   Non-Null Count  Dtype
```

```
---  ---
```

```
0    id      29345 non-null   int64
```

```
1    label   29345 non-null   int64
```

```
2    tweet   29345 non-null   object
```

```
dtypes: int64(2), object(1)
```

```
memory usage: 917.0+ KB
```

```
tweet_df['label'].value_counts()
```

```
label
```

```
0    27352
```


```
1     1993
```

```
Name: count, dtype: int64
```

```
fig = plt.figure(figsize=(5,5))
```

```
sns.countplot(x='label', data = tweet_df)
```


<Axes: xlabel='label', ylabel='count'>



```
fig = plt.figure(figsize=(7,7))
colors = ("red", "gold")
wp = {'linewidth':2, 'edgecolor':"black"}
tags = tweet_df['label'].value_counts()
explode = (0.1, 0.1)
tags.plot(kind='pie',autopct = '%1.1f%%', shadow=True, colors = colors, startangle =90,
          wedgeprops = wp, explode = explode, label='')
plt.title('Distribution of sentiments')
```

Text(0.5, 1.0, 'Distribution of sentiments')

Distribution of sentiments



```
non_hate_tweets = tweet_df[tweet_df.label == 0]
non_hate_tweets.head()
```

	id	label	tweet	
0	1	0	user father dysfunctional selfish drags kids d...	
1	2	0	user user thanks lyft credit cant use cause do...	
2	3	0	bihday majesty	
3	4	0	model love u take u time ur	
4	5	0	factsguide society motivation	

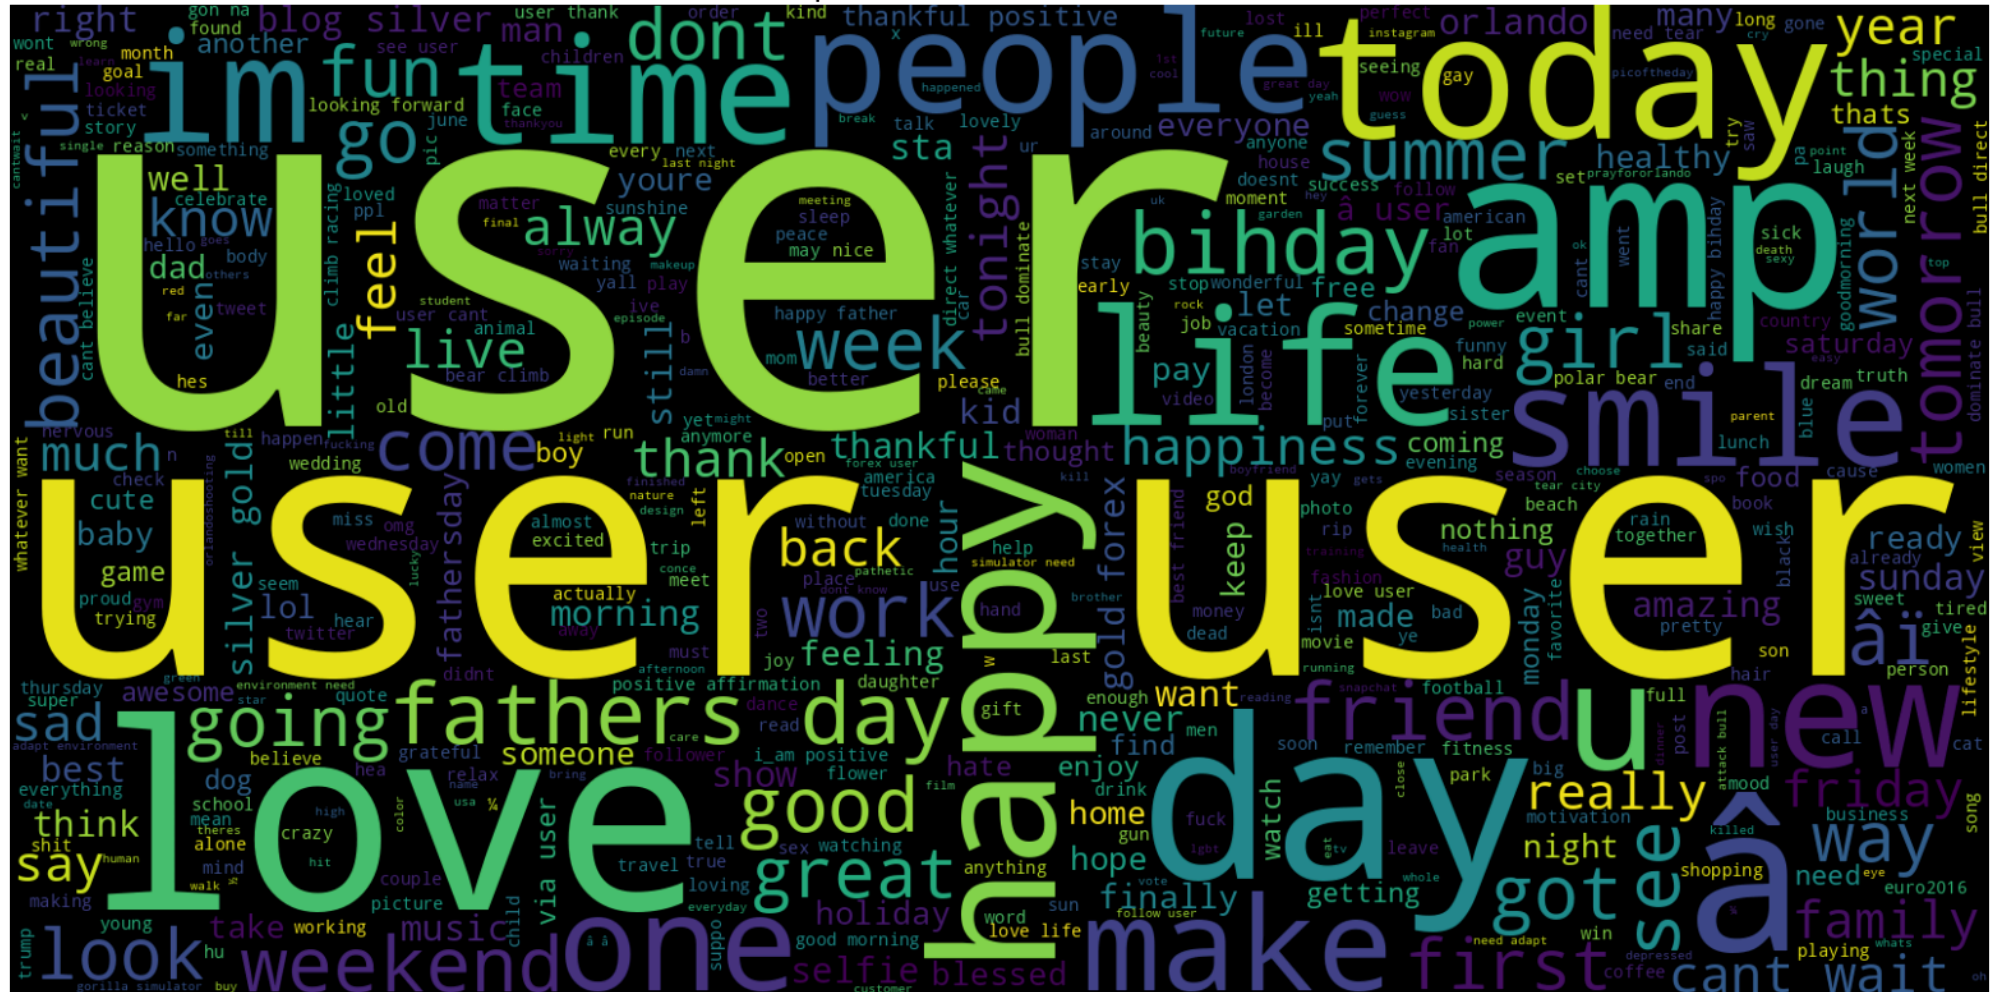
Next steps:

[Generate code with non_hate_tweets](#)



☒ [View recommended plots](#)

```
text = ' '.join([word for word in non_hate_tweets['tweet']])
plt.figure(figsize=(20,15), facecolor='None')
wordcloud = WordCloud(max_words=500, width=1600, height=800).generate(text)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Most frequent words in non hate tweets', fontsize = 19)
```

Most frequent words in non hate tweets



<https://colab.research.google.com/drive/1NcXjdadpLPY21dkLHFxdkgLT140HTfxW?authuser=5#scrollTo=cBLgZFJI0kkR&printMode=true>

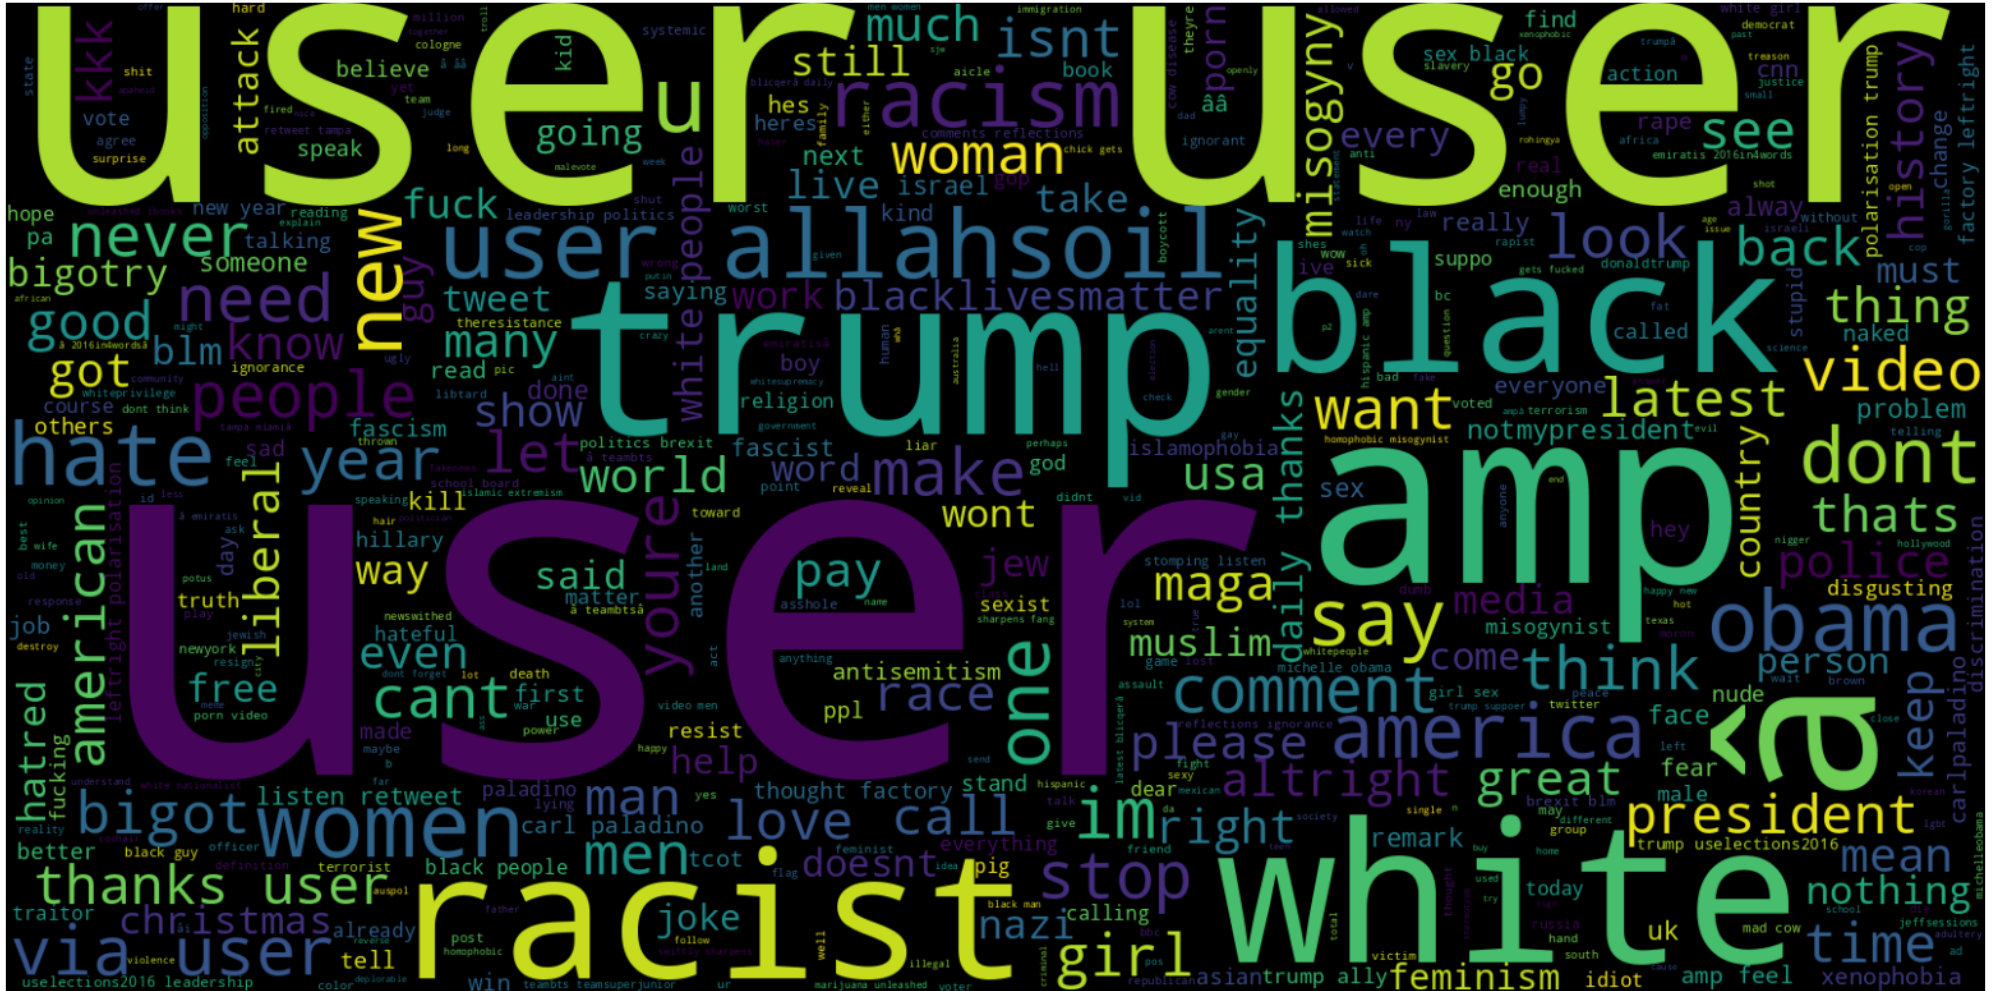
	id	label	tweet	
13	14	1	user cnn calls michigan middle school build wa...	
14	15	1	comment australia opkillingbay seashepherd hel...	
17	18	1	retweet agree	
23	24	1	user user lumpy says prove lumpy	
34	35	1	unbelievable 21st century wed need something l...	

Next steps:

[Generate code with neg_tweets](#) [View recommended plots](#)

```
text = ' '.join([word for word in neg_tweets['tweet']])
plt.figure(figsize=(20,15), facecolor='None')
wordcloud = WordCloud(max_words=500, width=1600, height=800).generate(text)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.title('Most frequent words in hate tweets', fontsize = 19)
plt.show()
```

Most frequent words in hate tweets



```
vect = TfidfVectorizer(ngram_range=(1,2)).fit(tweet_df['tweet'])
```

```
feature_names = vect.get_feature_names_out()
print("Number of features: {}".format(len(feature_names)))
print("First 20 features: \n{}".format(feature_names[:20]))
```

Number of features: 208579

```
First 20 features:
['0000001' '0000001 polluting' '00027' '00027 photooftheday' '001' '0035'
 '00h30' '01' '01 4995' '01 7900' '01 blog' '01 croatia' '01 may'
```



```

'01 shopalyssas' '0115' '0115 8599968' '0161' '0161 manny' '019'
'019 previous']

vect = TfidfVectorizer(ngram_range=(1,3)).fit(tweet_df['tweet'])

feature_names = vect.get_feature_names_out()
print("Number of features: {}".format(len(feature_names)))
print("First 20 features: \n{}".format(feature_names[:20]))

Number of features: 380305

First 20 features:
['0000001' '0000001 polluting' '0000001 polluting niger' '00027'
'00027 photooftheday' '00027 photooftheday music' '001' '0035' '00h30'
'01' '01 4995' '01 4995 rustic' '01 7900' '01 7900 shopalyssas' '01 blog'
'01 blog silver' '01 croatia' '01 croatia happy' '01 may' '01 may actual']

% Model building

X = tweet_df['tweet']
Y = tweet_df['label']
X = vect.transform(X)

x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, random_state=42)

print("Size of x_train:", (x_train.shape))
print("Size of y_train:", (y_train.shape))
print("Size of x_test: ", (x_test.shape))
print("Size of y_test: ", (y_test.shape))

Size of x_train: (23476, 380305)
Size of y_train: (23476,)
Size of x_test: (5869, 380305)
Size of y_test: (5869,)

logreg = LogisticRegression()
logreg.fit(x_train, y_train)
logreg_predict = logreg.predict(x_test)
logreg_acc = accuracy_score(logreg_predict, y_test)
print("Test accuracy: {:.2f}%".format(logreg_acc*100))

Test accuracy: 93.17%

print(confusion_matrix(y_test, logreg_predict))
print("\n")
print(classification_report(y_test, logreg_predict))

[[5458   0]
 [ 401  10]]

precision    recall  f1-score   support
```

	0	0.93	1.00	0.96	5458
	1	1.00	0.02	0.05	411
accuracy				0.93	5869
macro avg		0.97	0.51	0.51	5869
weighted avg		0.94	0.93	0.90	5869

```
style.use('classic')
cm = confusion_matrix(y_test, logreg_predict, labels=logreg.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=logreg.classes_)
disp.plot()
```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7ebe53661300>

